

Executive Profile

Engineering leader with **14+ years** building and running platform, data, and AI teams for connected-mobility and regulated enterprise products. Owned architecture and delivery of an AWS streaming lakehouse processing **5TB+/day across 500K+ connected vehicles**; now lead **Azure-native, multi-tenant GenAI platform engineering** for global enterprise tenants across India, UK, and US. Combines distributed-systems depth with people leadership — building pods, setting CI/CD and architecture governance standards, and running cloud cost governance as an engineering discipline. Targeting **Head of Platform Engineering, Director of Engineering, or Head of AI/Data Platform**.

5TB+ / day Streaming data processed	500K+ Connected vehicles on platform	Sub-second Analytics SLA at peak load
25–30% Cloud cost run-rate reduction	50% Fewer deployment failures	40% Faster release cycles

Leadership Scope

- **Teams:** Led pods of **8 engineers** directly, **mentored up to 12** across backend, data, and ML — hiring, onboarding, review standards, and career development.
- **Cross-functional:** Drove delivery across backend, data, AI, and security tracks with product, risk, compliance, and operations leaders.
- **Global governance:** Owned architecture direction and roadmap alignment with **India, UK, and US** stakeholders, including executive design reviews.
- **Standards & cost:** Institutionalized CI/CD, MLOps, and architecture governance (L1/L2/L3) — **50% fewer deployment failures, 40% faster releases, 25–30%** sustained cloud run-rate reduction.

Core Competencies

- **Platform Engineering:** Multi-tenant SaaS architecture · reusable platform IP · internal developer tooling · CI/CD standards · architecture governance (L1/L2/L3)
- **Data & Distributed Systems:** Event-driven ingestion · streaming lakehouse (Iceberg on S3) · warehouse modernization · exactly-once semantics · observability and data quality at scale
- **AI / GenAI Engineering:** LangGraph multi-agent orchestration · RAG and document intelligence · LLMOps and model governance · prompt-injection and guardrail design · zero-trust LLM brokering
- **Leadership & FinOps:** Team building · hiring and mentoring · roadmap ownership · cross-geo stakeholder management · design-time cost mapping and token economics

Technical Stack

- **Languages:** Python · SQL · Bash
- **AI & Agentic:** LangGraph · multi-agent systems · MCP gateways · RAG · Azure OpenAI · Azure AI Document Intelligence · zero-trust LLM brokering
- **Azure Platform:** Azure Functions · Cosmos DB · Microsoft Entra ID (OIDC, app roles) · VNet and Private Endpoints · Azure RBAC/ABAC · Key Vault
- **AWS Data & Lakehouse:** Apache Iceberg · S3 · Glue Streaming ETL · Kinesis · Redshift Serverless · SageMaker · Spark · Kafka · Hadoop · Databricks
- **Platform & Delivery:** Docker · Kubernetes · CI/CD · MLOps · multi-tenant isolation · SOC2/GDPR-aligned controls

Professional Experience

Associate Director of Engineering — Central AI & Platform Product Hub | KPMG India

Dec 2025 – Present | Bangalore, India

Scope: Architecture and engineering lead for a centralized product group building **Azure-native, multi-tenant software IP** for global enterprise tenants — platform strategy, GenAI security controls, and cross-team engineering standards.

TPRM — Multi-Tenant Enterprise Risk Platform

- Own end-to-end platform architecture for a cloud-native Third Party Risk Management product, standardizing questionnaire workflows, document review, and compliance decisioning into a reusable operating model.
- Replaced synchronous point-to-point integration with an **event-driven topology** to absorb long-running, callback-heavy risk processes — improving resiliency, horizontal scale, and audit traceability across tenants.
- Built intelligent document processing on **Azure AI Document Intelligence**, turning unstructured compliance evidence into queryable platform assets and cutting manual review cycles.
- Established architecture governance and security baselines across L1/L2/L3 flows — data privacy, prompt-injection defense, tenant boundaries — creating a reusable GenAI control plane for other product lines.

Architecture Copilot Studio — GenAI Architecture Intelligence Platform

- Lead a shared architecture intelligence platform converting prompts and Mermaid definitions into editable Draw.io artifacts, reducing dependence on manual architecture reviews.
- Replaced linear DAG prompt-chaining with **cyclic, stateful LangGraph multi-agent orchestration** over an append-only transaction log — recoverable execution, auditable reasoning traces, controlled iteration.
- Implemented a **zero-trust Decoupled Security Broker Pattern**: the LLM runtime operates as an untrusted container while privileged execution is offloaded to an **isolated Azure Functions broker inside a private VNet**, reached only over **Private Endpoints** — removing model-side data-plane access.
- Deployed **out-of-band authorization** pairing **Microsoft Entra ID OIDC JWT** validation with **app role and scope evaluation** for RBAC, plus tenant-claim cross-referencing against Cosmos DB ownership metadata for ABAC — preventing multi-tenant data leakage independent of model output.
- Hardened the broker with **circuit breaking and sanitized error contracts** so prompt-injection or parameter-tampering attempts fail closed with no data or system detail returned to the model.
- Introduced **design-time cost mapping** and scale-based cloud cost estimation at architecture inception, preventing recursive agent-loop token spikes before build commitment.

Principal Data Engineer — Platform Owner, Deep View Analytics | Darby Telematics India & Danlaw Inc

Mar 2019 – Nov 2025 | Bangalore, India

Scope: Platform owner for the connected-mobility data and AI backbone — **5TB+/day, 500K+ connected vehicles**. Led **8 engineers**, mentored up to **12** across data and ML, and owned platform cloud cost governance.

- Owned platform strategy for an AWS-native streaming lakehouse serving vehicle health, predictive maintenance, driver risk scoring, and emissions analytics — unifying ingestion, enrichment, ML, and self-service analytics on one governed backbone.
- Migrated provisioned **EMR/EC2 clusters to an event-driven serverless fabric** (Kinesis, Lambda, Glue Streaming ETL, Apache Iceberg on S3, Redshift Serverless) — **30% cloud cost reduction** while holding sub-second SLAs at 5TB+/day.
- Resolved network-dropout packet bursts with a **time-windowed lookback buffer guaranteeing exactly-once processing**, eliminating duplicate and lost fleet events across volatile connectivity.
- Institutionalized **SageMaker MLOps pipelines** for predictive maintenance, driver risk, and emissions forecasting — moving models from experimentation into governed production inference.
- Built the engineering operating model — CI/CD, data-quality automation, review standards, onboarding — driving **40% faster releases** and **50% fewer deployment failures** while growing pod capability.
- Owned FinOps across storage tiering, Iceberg compaction, and query patterns, sustaining **30% cost optimization** across half a million endpoints.

Senior Software Engineer — NEC India Pvt Ltd

Aug 2012 – Feb 2019 | Bangalore, India

- Architected Hadoop-based ingestion and processing pipelines for multi-source enterprise datasets, standardizing validation and transformation contracts for downstream reliability.
- Delivered a computer-vision Queue Length Management system (TensorFlow, OpenCV) for near-real-time retail congestion detection — **30% lower customer wait times**.
- Built a Predictive Auto Ordering System converting demand signals into inventory recommendations (**20% fewer stockouts**); shipped NEC's IoT Big Data Solution, cutting client setup time **50%**.

Education

- **M.Tech, Data Science & Engineering** — BITS Pilani (2020–2022)
- **B.Tech, Electronics & Communication** — GEIT, Dehradun (2007–2011)